

## ▼ Setup

### Install Dependencies

```
!pip install gymnasium
```

```
Requirement already satisfied: gymnasium in /usr/local/lib/python3.10.0/dist-packages
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10.0/dist-packages
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.10.0/dist-packages
Requirement already satisfied: typing-extensions>=4.1.1 in /usr/local/lib/python3.10.0/dist-packages
Requirement already satisfied: farama-notifications>=0.0.4 in /usr/local/lib/python3.10.0/dist-packages
```

### Import dependencies

```
import gymnasium as gym
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.distributions import Categorical
import time
import matplotlib.pyplot as plt
```

## ▼ Preperation

### Define the Policy Network

```
class Policy(nn.Module):
    def __init__(self, state_dim, action_dim):
        super(Policy, self).__init__()
        self.fc1 = nn.Linear(state_dim, 128)
        self.fc2 = nn.Linear(128, action_dim)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return F.softmax(x, dim=1)
```

Resources

app.log X

...

1 to 10 of 60 entries

Filter



Timestamp	Level	Message
Jul 14, 2023, 11:34:34 PM	WARNING	WARNING:root:3e91514d-3ca4-9804-8c1f3a2a restarted
Jul 14, 2023, 11:34:34 PM	INFO	KernelRestart: restarting kernel keep random p
Jul 14, 2023, 11:34:33 PM	INFO	Kernel interrupt: 3e91514d-3ca4-9804-8c1f3a2a
Jul 14, 2023, 11:34:32 PM	INFO	Kernel interrupt: 3e91514d-3ca4-9804-8c1f3a2a
Jul 14, 2023, 11:34:26 PM	INFO	Kernel interrupt: 3e91514d-3ca4-9804-8c1f3a2a
Jul 14, 2023, 10:16:07 PM	INFO	Kernel started: 3e91514d-3ca4-9804-8c1f3a2a name: python3
Jul 14, 2023, 10:11:46 PM	INFO	Use Control-C this server and down all kernel to skip confirm
Jul 14, 2023, 10:11:46 PM	INFO	http://172.28.0.
Jul 14, 2023, 10:11:46 PM	INFO	Jupyter Notebo is running at:
Jul 14, 2023, 10:11:46 PM	INFO	Serving notebo local directory:

Show 10 per page

1

2

3

4

5

6

Create the environment, instantiate the policy network and define the optimizer

```
# Create the environment
env = gym.make('CartPole-v1',render_mode="rgb_array")
state_dim = env.observation_space.shape[0]
action_dim = env.action_space.n

# Initialize the policy network
policy = Policy(state_dim, action_dim)

# Define the optimizer
optimizer = optim.Adam(policy.parameters(), lr=0.01)
```

## ▼ Algorithm

Pick an action based on policy

```
def select_action(state):
    state = np.array(state)
    state = torch.from_numpy(state).float().unsqueeze(0)
    probs = policy(state)
    m = Categorical(probs)
    action = m.sample()
    return action.item(), m.log_prob(action)
```

Policy Gradient Algorithm, the actual training loop

```
def policy_gradient():
    num_episodes = 1000
    gamma = 0.99

    rewards_per_episode = [] # List to store rewards for

    # for 1000 episodes
    for episode in range(num_episodes):
        observations = env.reset()
        state = np.array(observations[0])
        episode_reward = 0
        log_probs = []
        rewards = []

        # loop through each time step in one episode
        while True:
```

```

    action, log_prob = select_action(state)
    next_state, reward, done, _, _ = env.step(action)

    log_probs.append(log_prob)
    rewards.append(reward)
    episode_reward += reward

    if done:
        break

    state = next_state

# Compute the discounted rewards
discounts = [gamma**i for i in range(len(rewards))]
discounted_rewards = [discount * reward for discount, reward in zip(discounts, rewards)]

# Convert the discounted_rewards into a Tensor
discounted_rewards = torch.Tensor(discounted_rewards)

# Normalize the discounted rewards
discounted_rewards -= torch.mean(discounted_rewards)
discounted_rewards /= torch.std(discounted_rewards)

# Calculate the loss
policy_loss = []
for log_prob, reward in zip(log_probs, discounted_rewards):
    policy_loss.append(-log_prob * reward)
policy_loss = torch.cat(policy_loss).sum()

# Update the policy network
optimizer.zero_grad()
policy_loss.backward()
optimizer.step()

# Print the episode statistics
if episode % 10 == 0:
    print('Episode {}: reward = {}'.format(episode, episode_reward))

# Plot the rewards per episode
plt.plot(rewards_per_episode)
plt.xlabel('Episode')
plt.ylabel('Reward')
plt.title('Reward per Episode')
plt.show()

```

## ▼ Run Trials

```
policy_gradient()
```

```
Episode 0: reward = 11.0
Episode 10: reward = 31.0
Episode 20: reward = 11.0
Episode 30: reward = 72.0
Episode 40: reward = 31.0
Episode 50: reward = 124.0
Episode 60: reward = 50.0
Episode 70: reward = 36.0
Episode 80: reward = 42.0
Episode 90: reward = 203.0
Episode 100: reward = 112.0
Episode 110: reward = 53.0
Episode 120: reward = 106.0
Episode 130: reward = 93.0
Episode 140: reward = 65.0
Episode 150: reward = 41.0
Episode 160: reward = 87.0
Episode 170: reward = 97.0
Episode 180: reward = 44.0
Episode 190: reward = 62.0
Episode 200: reward = 114.0
Episode 210: reward = 70.0
Episode 220: reward = 83.0
Episode 230: reward = 51.0
Episode 240: reward = 49.0
Episode 250: reward = 39.0
Episode 260: reward = 89.0
Episode 270: reward = 70.0
Episode 280: reward = 46.0
Episode 290: reward = 36.0
Episode 300: reward = 46.0
Episode 310: reward = 61.0
Episode 320: reward = 68.0
Episode 330: reward = 110.0
Episode 340: reward = 135.0
Episode 350: reward = 134.0
Episode 360: reward = 175.0
Episode 370: reward = 115.0
Episode 380: reward = 78.0
Episode 390: reward = 160.0
Episode 400: reward = 388.0
Episode 410: reward = 144.0
Episode 420: reward = 98.0
Episode 430: reward = 63.0
Episode 440: reward = 58.0
Episode 450: reward = 47.0
Episode 460: reward = 64.0
Episode 470: reward = 119.0
Episode 480: reward = 114.0
Episode 490: reward = 115.0
Episode 500: reward = 102.0
Episode 510: reward = 114.0
Episode 520: reward = 63.0
Episode 530: reward = 110.0
```

Episode 540: reward = 121.0

Episode 550: reward = 114.0

Episode 560: reward = 138.0

Episode 570: reward = 135.0

[Colab paid products](#) - [Cancel contracts here](#)



54m 16s completed at 11:34 PM

